

---

# **PGS logistic regression optimization**

**Razvan Sultana**

**Nov 09, 2022**



# CONTENTS

<b>1</b>	<b>Cooper 142 SNPs set</b>	<b>3</b>
1.1	Preparation . . . . .	3
1.2	All participants . . . . .	5
1.3	Males . . . . .	7
1.4	Females . . . . .	9
1.5	NN Males . . . . .	11
1.6	NI Males . . . . .	13
1.7	II Males . . . . .	15
1.8	NN Females . . . . .	17
1.9	NI Females . . . . .	20
1.10	II Females . . . . .	22



This project applies logistic regression (LR) to rank SNPs from several PGS datasets as predictors of Parkinson's disease. Click on each dataset link to explore the details.

- [\*Cooper 142 SNPs set\*](#)



## COOPER 142 SNPS SET

### 1.1 Preparation

#### 1.1.1 Import required packages.

```
import os, sys, warnings
import numpy as np
import pandas as pd
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import GridSearchCV
from sklearn.exceptions import ConvergenceWarning
```

#### 1.1.2 Read input matrix with genotypes

The matrix contains the genotypes from AMP-PD/MGRB dataset for 140 SNPs.

```
table = pd.read_csv("data/matrix.txt", sep="\t")
table
```

	participant_id	phenotype	cohort	gender	inv_genotype	\
0	SY-NIH_INVAA791MKCET	1	STEADY-PD3	M	NI	
1	SY-NIH_INVFP886EEYYL	1	STEADY-PD3	M	NI	
2	SY-NIH_INVFM717GWDX4	1	STEADY-PD3	F	NI	
3	SY-NIH_INVNN611MKKN9	1	STEADY-PD3	M	NI	
4	SY-NIH_INVRB171EXGUK	1	STEADY-PD3	M	II	
...	...	...	...	...	...	...
3107	BABQX	0	MGRB	M	II	
3108	BABRB	0	MGRB	F	II	
3109	BABRE	0	MGRB	M	NI	
3110	ZAAAB	0	MGRB	M	NI	
3111	AABUO	0	MGRB	F	NN	
	rs2275579	rs144115304	rs115581042	rs79531911	rs138844738	... \
0	0	0	0	0	0	...
1	0	0	0	0	0	...
2	0	0	0	0	0	...
3	0	0	0	0	0	...
4	0	1	1	1	0	...

(continues on next page)

(continued from previous page)

```

...
3107      0      0      0      0      0      0 ...
3108      0      0      0      0      0      0 ...
3109      0      0      0      0      0      0 ...
3110      0      0      0      0      0      0 ...
3111      0      0      0      0      0      0 ...

      rs10448130  rs34288580  rs34992950  rs7387252  rs2410595  rs41311559  \
0           2           2           1           2           1           0
1           0           0           0           0           0           0
2           0           0           0           0           0           0
3           1           0           1           1           0           0
4           1           1           0           1           0           0
...
3107      1           0           0           0           0           0
3108      0           0           0           0           0           0
3109      1           0           0           0           0           0
3110      2           1           1           1           0           0
3111      1           0           0           0           0           0

      rs148894916  rs112957100  rs143756122  rs148514732
0           0           0           0           0
1           0           0           0           0
2           0           0           0           0
3           0           0           0           0
4           0           0           0           0
...
3107      0           0           0           0
3108      0           0           0           0
3109      0           0           0           0
3110      0           0           0           0
3111      0           0           0           0

[3112 rows x 145 columns]

```

### 1.1.3 Distribution of data

#### Distribution by phenotype

(0=Control, 1=Case)

```
table.groupby('phenotype')['participant_id'].nunique()
```

```

phenotype
0      1556
1      1556
Name: participant_id, dtype: int64

```



## Distribution by gender/phenotype

```
table.groupby(['gender', 'phenotype'])['participant_id'].nunique()
```

```
gender  phenotype
F        0           567
         1           567
M        0          989
         1          989
Name: participant_id, dtype: int64
```

## Distribution by gender/phenotype/inv8\_001 genotype

```
table.groupby(['gender', 'phenotype', 'inv_genotype'])['participant_id'].nunique()
```

```
gender  phenotype  inv_genotype
F        0          II           195
          NI           259
          NN           113
         1          II           175
          NI           270
          NN           122
M        0          II           318
          NI           480
          NN           191
         1          II           296
          NI           477
          NN           216
Name: participant_id, dtype: int64
```

## 1.2 All participants

### 1.2.1 Logistic regression model

```
pd.set_option('display.max_rows', 150)
X = table[table.columns[5:]]
Y = table['phenotype']
lr = LogisticRegression(random_state=42, solver='saga', n_jobs=-1, penalty='elasticnet
↳')
table.groupby('phenotype')['participant_id'].nunique()
```

```
phenotype
0      1556
1      1556
Name: participant_id, dtype: int64
```

## 1.2.2 Grid search for 3 hyperparameters

```
# parameters = {'C': [0.005, 0.01, 0.02, 0.05, 0.1, 0.5, 1, 10, 20, 30],
#               'max_iter': [10, 25, 50, 75, 100, 150, 200, 400, 800, 1600],
#               'l1_ratio': [1, 0.9, 0.8, 0.7, 0.6, 0.5, 0.4, 0.3, 0.2, 0.1]}

parameters = {'C': [0.01, 0.02, 0.05],
               'max_iter': [10, 25, 50],
               'l1_ratio': [1, 0.9, 0.8]}

grid_lr = GridSearchCV(lr, parameters, verbose=False, scoring='roc_auc', n_jobs=-1,
cv=10)
if not sys.warnoptions:
    warnings.simplefilter("ignore", category=ConvergenceWarning)
    os.environ["PYTHONWARNINGS"] = "ignore" # Also affect subprocesses
grid_lr.fit(X, Y)
```

```
GridSearchCV(cv=10,
             estimator=LogisticRegression(n_jobs=-1, penalty='elasticnet',
                                           random_state=42, solver='saga'),
             n_jobs=-1,
             param_grid={'C': [0.01, 0.02, 0.05], 'l1_ratio': [1, 0.9, 0.8],
                          'max_iter': [10, 25, 50]},
             scoring='roc_auc', verbose=False)
```

## 1.2.3 Best estimator

```
best_lr = grid_lr.best_estimator_

max_auc_score = roc_auc_score(Y, best_lr.predict_proba(X)[:, 1])
coefs = best_lr.coef_[0, :]
num_coef = np.sum(coefs != 0)
X_header = np.array(X.columns)

data_array = np.vstack((X_header, coefs))
model_coefs = pd.DataFrame(data=data_array.T, columns=['SNP', 'Coefficient'])
print(f'Max AUC score: {max_auc_score}\n')
print(f'Non-zero coefficients: {num_coef}\n')
print(f'Best estimator: {grid_lr.best_estimator_}')
print(f'Scorer: {grid_lr.scorer_}')
print(f'Best params: {grid_lr.best_params_}')
print(f'Best score: {grid_lr.best_score_}\n')
m = model_coefs[model_coefs['Coefficient'] != 0].sort_values(by='Coefficient')
m = m.reset_index(drop=True).assign(Index=range(len(m)))
m.Index = m.Index + 1
m.set_index('Index')
```

Max AUC score:0.5551641873897212

Non-zero coefficients: 2

Best estimator: LogisticRegression(C=0.02, l1\_ratio=1, max\_iter=25, n\_jobs=-1,
penalty='elasticnet', random\_state=42, solver='saga')

(continues on next page)

(continued from previous page)

```
Scorer: make_scorer(roc_auc_score, needs_threshold=True)
Best params: {'C': 0.02, 'l1_ratio': 1, 'max_iter': 25}
Best score: 0.5492252417764205
```

Index		SNP Coefficient
1	rs11248057	0.055163
2	rs3806760	0.098122

## 1.3 Males

### 1.3.1 Logistic regression model

```
table1 = table[table.gender == "M"]
X = table1[table1.columns[5:]]
Y = table1['phenotype']
lr = LogisticRegression(random_state=42, solver='saga', n_jobs=-1, penalty='elasticnet',
                        cv=10)
table1.groupby('phenotype')['participant_id'].nunique()
```

```
phenotype
0      989
1      989
Name: participant_id, dtype: int64
```

### 1.3.2 Grid search for 3 hyperparameters

```
# parameters = {'C': [0.005, 0.01, 0.02, 0.05, 0.1, 0.5, 1, 10, 20, 30],
#               'max_iter': [10, 25, 50, 75, 100, 150, 200, 400, 800, 1600],
#               'l1_ratio': [1, 0.9, 0.8, 0.7, 0.6, 0.5, 0.4, 0.3, 0.2, 0.1]}

parameters = {'C': [0.01, 0.02],
               'max_iter': [75, 100, 150],
               'l1_ratio': [0.2, 0.1]}

grid_lr = GridSearchCV(lr, parameters, verbose=False, scoring='roc_auc', n_jobs=-1,
                       cv=10)
if not sys.warnoptions:
    warnings.simplefilter("ignore", category=ConvergenceWarning)
    os.environ["PYTHONWARNINGS"] = "ignore" # Also affect subprocesses
grid_lr.fit(X, Y)
```

```
GridSearchCV(cv=10,
             estimator=LogisticRegression(n_jobs=-1, penalty='elasticnet',
                                           random_state=42, solver='saga'),
             n_jobs=-1,
             param_grid={'C': [0.01, 0.02], 'l1_ratio': [0.2, 0.1]},
```

(continues on next page)

(continued from previous page)

```

        'max_iter': [75, 100, 150]},
        scoring='roc_auc', verbose=False)

```

### 1.3.3 Best estimator

```

best_lr = grid_lr.best_estimator_

max_auc_score = roc_auc_score(Y, best_lr.predict_proba(X)[:, 1])
coefs = best_lr.coef_[0, :]
num_coef = np.sum(coefs != 0)
X_header = np.array(X.columns)

data_array = np.vstack((X_header, coefs))
model_coefs = pd.DataFrame(data=data_array.T, columns=['SNP', 'Coefficient'])
print(f'Max AUC score:{max_auc_score}\n')
print(f'Non-zero coefficients: {num_coef}\n')
print(f'Best estimator: {grid_lr.best_estimator_}')
print(f'Scorer: {grid_lr.scorer_}')
print(f'Best params: {grid_lr.best_params_}')
print(f'Best score: {grid_lr.best_score_}\n')
m = model_coefs[model_coefs['Coefficient'] != 0].sort_values(by='Coefficient')
m = m.reset_index(drop=True).assign(Index=range(len(m)))
m.Index= m.Index + 1
m.set_index('Index')

```

Max AUC score:0.6243225531401534

Non-zero coefficients: 37

Best estimator: LogisticRegression(C=0.01, l1\_ratio=0.1, n\_jobs=-1, penalty=  
↪'elasticnet',

random\_state=42, solver='saga')

Scorer: make\_scorer(roc\_auc\_score, needs\_threshold=True)

Best params: {'C': 0.01, 'l1\_ratio': 0.1, 'max\_iter': 100}

Best score: 0.5620659282996945

Index	SNP	Coefficient
1	rs7520918	-0.10687
2	rs112270735	-0.077184
3	rs4444618	-0.054004
4	rs13161496	-0.044109
5	rs3733349	-0.035941
6	rs3785891	-0.022054
7	rs494312	-0.015573
8	rs62190394	-0.008948
9	rs12995314	-0.00536
10	rs72838312	-0.004366
11	rs142448570	-0.00037
12	rs74609071	0.000361
13	rs4028634	0.0029
14	rs17016235	0.00292

(continues on next page)

(continued from previous page)

15	rs1501467	0.006768
16	rs7682766	0.007628
17	rs6532190	0.010966
18	rs379066	0.015095
19	rs2412116	0.015274
20	rs5848	0.015471
21	rs764324	0.01817
22	rs11601088	0.019219
23	rs10110312	0.028269
24	rs4897753	0.029363
25	rs6599389	0.03295
26	rs11097213	0.036052
27	rs3822023	0.041542
28	rs79828056	0.046307
29	rs1372420	0.051416
30	rs11264302	0.057684
31	rs10831599	0.065901
32	rs12434297	0.068324
33	rs9985581	0.070506
34	rs6842271	0.077488
35	rs12935995	0.079672
36	rs11248057	0.080207
37	rs3806760	0.147021

## 1.4 Females

### 1.4.1 Logistic regression model

```
table1 = table[table.gender == "F"]
X = table1[table1.columns[5:]]
Y = table1['phenotype']
lr = LogisticRegression(random_state=42, solver='saga', n_jobs=-1, penalty='elasticnet
↳')
table1.groupby('phenotype')['participant_id'].nunique()
```

```
phenotype
0      567
1      567
Name: participant_id, dtype: int64
```

### 1.4.2 Grid search for 3 hyperparameters

```
# parameters = {'C': [0.005, 0.01, 0.02, 0.05, 0.1, 0.5, 1, 10, 20, 30],
#               'max_iter': [10, 25, 50, 75, 100, 150, 200, 400, 800, 1600],
#               'l1_ratio': [1, 0.9, 0.8, 0.7, 0.6, 0.5, 0.4, 0.3, 0.2, 0.1]}

parameters = {'C': [0.005, 0.01],
               'max_iter': [25, 50, 75],
               'l1_ratio': [0.1, 0.2]}
```

(continues on next page)

(continued from previous page)

```

grid_lr = GridSearchCV(lr, parameters, verbose=False, scoring='roc_auc', n_jobs=-1,
cv=10)
if not sys.warnoptions:
    warnings.simplefilter("ignore", category=ConvergenceWarning)
    os.environ["PYTHONWARNINGS"] = "ignore" # Also affect subprocesses
grid_lr.fit(X, Y)

```

```

GridSearchCV(cv=10,
             estimator=LogisticRegression(n_jobs=-1, penalty='elasticnet',
                                           random_state=42, solver='saga'),
             n_jobs=-1,
             param_grid={'C': [0.005, 0.01], 'l1_ratio': [0.1, 0.2],
                          'max_iter': [25, 50, 75]},
             scoring='roc_auc', verbose=False)

```

### 1.4.3 Best estimator

```

best_lr = grid_lr.best_estimator_

max_auc_score = roc_auc_score(Y, best_lr.predict_proba(X)[:, 1])
coefs = best_lr.coef_[0, :]
num_coef = np.sum(coefs != 0)
X_header = np.array(X.columns)

data_array = np.vstack((X_header, coefs))
model_coefs = pd.DataFrame(data=data_array.T, columns=['SNP', 'Coefficient'])
print(f'Max AUC score: {max_auc_score}\n')
print(f'Non-zero coefficients: {num_coef}\n')
print(f'Best estimator: {grid_lr.best_estimator_}')
print(f'Scorer: {grid_lr.scorer_}')
print(f'Best params: {grid_lr.best_params_}')
print(f'Best score: {grid_lr.best_score_}\n')
m = model_coefs[model_coefs['Coefficient'] != 0].sort_values(by='Coefficient')
m = m.reset_index(drop=True).assign(Index=range(len(m)))
m.Index = m.Index + 1
m.set_index('Index')

```

Max AUC score: 0.56756840825036

Non-zero coefficients: 6

Best estimator: LogisticRegression(C=0.005, l1\_ratio=0.1, max\_iter=50, n\_jobs=-1,
penalty='elasticnet', random\_state=42, solver='saga')

Scorer: make\_scorer(roc\_auc\_score, needs\_threshold=True)

Best params: {'C': 0.005, 'l1\_ratio': 0.1, 'max\_iter': 50}

Best score: 0.5115738908675196

	SNP	Coefficient
Index		
1	rs6871718	-0.026889

(continues on next page)

(continued from previous page)

2	rs1949362	-0.023602
3	rs6964	0.005616
4	rs11248057	0.005966
5	rs3850379	0.008878
6	rs3806760	0.016905

## 1.5 NN Males

### 1.5.1 Logistic regression model

```
table1 = table[(table.gender == "M") & (table.inv_genotype=="NN")]
X = table1[table1.columns[5:]]
Y = table1['phenotype']
lr = LogisticRegression(random_state=42, solver='saga', n_jobs=-1, penalty='elasticnet',
↳)
table1.groupby('phenotype')['participant_id'].nunique()
```

```
phenotype
0    191
1    216
Name: participant_id, dtype: int64
```

### 1.5.2 Grid search for 3 hyperparameters

```
# parameters = {'C': [0.005, 0.01, 0.02, 0.05, 0.1, 0.5, 1, 10, 20, 30],
#               'max_iter': [10, 25, 50, 75, 100, 150, 200, 400, 800, 1600],
#               'l1_ratio': [1, 0.9, 0.8, 0.7, 0.6, 0.5, 0.4, 0.3, 0.2, 0.1]}

parameters = {'C': [0.05, 0.1, 0.5],
               'max_iter': [10, 25],
               'l1_ratio': [0.7, 0.6, 0.5]}

grid_lr = GridSearchCV(lr, parameters, verbose=False, scoring='roc_auc', n_jobs=-1,
↳cv=10)
if not sys.warnoptions:
    warnings.simplefilter("ignore", category=ConvergenceWarning)
    os.environ["PYTHONWARNINGS"] = "ignore" # Also affect subprocesses
grid_lr.fit(X, Y)
```

```
GridSearchCV(cv=10,
             estimator=LogisticRegression(n_jobs=-1, penalty='elasticnet',
                                           random_state=42, solver='saga'),
             n_jobs=-1,
             param_grid={'C': [0.05, 0.1, 0.5], 'l1_ratio': [0.7, 0.6, 0.5],
                          'max_iter': [10, 25]},
             scoring='roc_auc', verbose=False)
```

### 1.5.3 Best estimator

```
best_lr = grid_lr.best_estimator_

max_auc_score = roc_auc_score(Y, best_lr.predict_proba(X)[:, 1])
coefs = best_lr.coef_[0, :]
num_coef = np.sum(coefs != 0)
X_header = np.array(X.columns)

data_array = np.vstack((X_header, coefs))
model_coefs = pd.DataFrame(data=data_array.T, columns=['SNP', 'Coefficient'])
print(f'Max AUC score: {max_auc_score}\n')
print(f'Non-zero coefficients: {num_coef}\n')
print(f'Best estimator: {grid_lr.best_estimator_}')
print(f'Scorer: {grid_lr.scorer_}')
print(f'Best params: {grid_lr.best_params_}')
print(f'Best score: {grid_lr.best_score_}\n')
m = model_coefs[model_coefs['Coefficient'] != 0].sort_values(by='Coefficient')
m = m.reset_index(drop=True).assign(Index=range(len(m)))
m.Index = m.Index + 1
m.set_index('Index')
```

Max AUC score: 0.7224403723094822

Non-zero coefficients: 32

Best estimator: LogisticRegression(C=0.1, l1\_ratio=0.6, max\_iter=10, n\_jobs=-1,  
penalty='elasticnet', random\_state=42, solver='saga')

Scorer: make\_scorer(roc\_auc\_score, needs\_threshold=True)

Best params: {'C': 0.1, 'l1\_ratio': 0.6, 'max\_iter': 10}

Best score: 0.5580189109136477

Index	SNP	Coefficient
1	rs35776335	-0.213769
2	rs13161496	-0.193039
3	rs12499663	-0.150666
4	rs444618	-0.124434
5	rs112270735	-0.083023
6	rs1269243287	-0.055458
7	rs3912643	-0.048141
8	rs11097297	-0.047875
9	rs2128786	-0.03413
10	rs3850379	-0.030484
11	rs4331494	-0.014515
12	rs11264302	-0.001455
13	rs3850745	-0.001154
14	rs35456861	-0.000152
15	rs764324	-0.000024
16	rs12543164	0.000034
17	rs10448130	0.000146
18	rs35933728	0.000247
19	rs2412116	0.018849
20	rs79828056	0.022762
21	rs1949362	0.027434
22	rs1501467	0.035676

(continues on next page)



(continued from previous page)

23	rs78197677	0.05128
24	rs1372420	0.073048
25	rs6842271	0.088079
26	rs12434297	0.089281
27	rs9985581	0.124415
28	rs3822023	0.128854
29	rs55805734	0.133538
30	rs4897753	0.210087
31	rs6964	0.215979
32	rs3806760	0.233441

## 1.6 NI Males

### 1.6.1 Logistic regression model

```
table1 = table[(table.gender == "M") & (table.inv_genotype=="NI")]
X = table1[table1.columns[5:]]
Y = table1['phenotype']
lr = LogisticRegression(random_state=42, solver='saga', n_jobs=-1, penalty='elasticnet',
↳ ')
table1.groupby('phenotype')['participant_id'].nunique()
```

```
phenotype
0      480
1      477
Name: participant_id, dtype: int64
```

### 1.6.2 Grid search for 3 hyperparameters

```
# parameters = {'C': [0.005, 0.01, 0.02, 0.05, 0.1, 0.5, 1, 10, 20, 30],
#               'max_iter': [10, 25, 50, 75, 100, 150, 200, 400, 800, 1600],
#               'l1_ratio': [1, 0.9, 0.8, 0.7, 0.6, 0.5, 0.4, 0.3, 0.2, 0.1]}

parameters = {'C': [0.005, 0.01],
               'max_iter': [10, 25, 50],
               'l1_ratio': [0.2, 0.1]}

grid_lr = GridSearchCV(lr, parameters, verbose=False, scoring='roc_auc', n_jobs=-1,
↳ cv=10)
if not sys.warnoptions:
    warnings.simplefilter("ignore", category=ConvergenceWarning)
    os.environ["PYTHONWARNINGS"] = "ignore" # Also affect subprocesses
grid_lr.fit(X, Y)
```

```
GridSearchCV(cv=10,
             estimator=LogisticRegression(n_jobs=-1, penalty='elasticnet',
                                           random_state=42, solver='saga'),
             n_jobs=-1,
```

(continues on next page)

(continued from previous page)

```
param_grid={'C': [0.005, 0.01], 'l1_ratio': [0.2, 0.1],
            'max_iter': [10, 25, 50]},
scoring='roc_auc', verbose=False)
```

### 1.6.3 Best estimator

```
best_lr = grid_lr.best_estimator_

max_auc_score = roc_auc_score(Y, best_lr.predict_proba(X)[:, 1])
coefs = best_lr.coef_[0, :]
num_coef = np.sum(coefs != 0)
X_header = np.array(X.columns)

data_array = np.vstack((X_header, coefs))
model_coefs = pd.DataFrame(data=data_array.T, columns=['SNP', 'Coefficient'])
print(f'Max AUC score: {max_auc_score}\n')
print(f'Non-zero coefficients: {num_coef}\n')
print(f'Best estimator: {grid_lr.best_estimator_}')
print(f'Scorer: {grid_lr.scorer_}')
print(f'Best params: {grid_lr.best_params_}')
print(f'Best score: {grid_lr.best_score_}\n')
m = model_coefs[model_coefs['Coefficient'] != 0].sort_values(by='Coefficient')
m = m.reset_index(drop=True).assign(Index=range(len(m)))
m.Index = m.Index + 1
m.set_index('Index')
```

Max AUC score: 0.5835888364779874

Non-zero coefficients: 4

Best estimator: LogisticRegression(C=0.005, l1\_ratio=0.1, max\_iter=25, n\_jobs=-1,  
penalty='elasticnet', random\_state=42, solver='saga')

Scorer: make\_scorer(roc\_auc\_score, needs\_threshold=True)

Best params: {'C': 0.005, 'l1\_ratio': 0.1, 'max\_iter': 25}

Best score: 0.5212733636229314

Index	SNP	Coefficient
1	rs7520918	-0.011899
2	rs9985581	0.016555
3	rs6842271	0.016555
4	rs3806760	0.027075

## 1.7 II Males

### 1.7.1 Logistic regression model

```
table1 = table[(table.gender == "M") & (table.inv_genotype=="II")]
X = table1[table1.columns[5:]]
Y = table1['phenotype']
lr = LogisticRegression(random_state=42, solver='saga', n_jobs=-1, penalty='elasticnet',
↳)
table1.groupby('phenotype')['participant_id'].nunique()
```

```
phenotype
0      318
1      296
Name: participant_id, dtype: int64
```

### 1.7.2 Grid search for 3 hyperparameters

```
# parameters = {'C': [0.005, 0.01, 0.02, 0.05, 0.1, 0.5, 1, 10, 20, 30],
#               'max_iter': [10, 25, 50, 75, 100, 150, 200, 400, 800, 1600],
#               'l1_ratio': [1, 0.9, 0.8, 0.7, 0.6, 0.5, 0.4, 0.3, 0.2, 0.1]}

parameters = {'C': [0.01, 0.02],
               'max_iter': [50, 75, 100],
               'l1_ratio': [0.2, 0.1]}

grid_lr = GridSearchCV(lr, parameters, verbose=False, scoring='roc_auc', n_jobs=-1,
↳cv=10)
if not sys.warnoptions:
    warnings.simplefilter("ignore", category=ConvergenceWarning)
    os.environ["PYTHONWARNINGS"] = "ignore" # Also affect subprocesses
grid_lr.fit(X, Y)
```

```
GridSearchCV(cv=10,
             estimator=LogisticRegression(n_jobs=-1, penalty='elasticnet',
                                           random_state=42, solver='saga'),
             n_jobs=-1,
             param_grid={'C': [0.01, 0.02], 'l1_ratio': [0.2, 0.1],
                          'max_iter': [50, 75, 100]},
             scoring='roc_auc', verbose=False)
```

### 1.7.3 Best estimator

```
best_lr = grid_lr.best_estimator_

max_auc_score = roc_auc_score(Y, best_lr.predict_proba(X)[:, 1])
coefs = best_lr.coef_[0, :]
num_coef = np.sum(coefs != 0)
X_header = np.array(X.columns)

data_array = np.vstack((X_header, coefs))
model_coefs = pd.DataFrame(data=data_array.T, columns=['SNP', 'Coefficient'])
print(f'Max AUC score:{max_auc_score}\n')
print(f'Non-zero coefficients: {num_coef}\n')
print(f'Best estimator: {grid_lr.best_estimator_}')
print(f'Scorer: {grid_lr.scorer_}')
print(f'Best params: {grid_lr.best_params_}')
print(f'Best score: {grid_lr.best_score_}\n')
m = model_coefs[model_coefs['Coefficient'] != 0].sort_values(by='Coefficient')
m = m.reset_index(drop=True).assign(Index=range(len(m)))
m.Index= m.Index + 1
m.set_index('Index')
```

Max AUC score:0.6759093999660037

Non-zero coefficients: 18

Best estimator: LogisticRegression(C=0.01, l1\_ratio=0.1, max\_iter=75, n\_jobs=-1,  
penalty='elasticnet', random\_state=42, solver='saga')

Scorer: make\_scorer(roc\_auc\_score, needs\_threshold=True)

Best params: {'C': 0.01, 'l1\_ratio': 0.1, 'max\_iter': 75}

Best score: 0.6007288190582128

Index	SNP	Coefficient
1	rs34333	-0.092096
2	rs3733349	-0.053093
3	rs3785891	-0.013793
4	rs8064765	-0.003056
5	rs62075803	-0.002828
6	rs4796663	-0.002828
7	rs4028634	0.001729
8	rs2410595	0.003947
9	rs764324	0.00623
10	rs4482120	0.009129
11	rs112318363	0.037641
12	rs12935995	0.040157
13	rs11097213	0.049525
14	rs3806760	0.050662
15	rs11248057	0.078256
16	rs11264302	0.08048
17	rs5848	0.085027
18	rs12434297	0.085108

## 1.8 NN Females

### 1.8.1 Logistic regression model

```
table1 = table[table.gender == "F" & (table.inv_genotype=="NN")]
X = table1[table1.columns[5:]]
Y = table1['phenotype']
lr = LogisticRegression(random_state=42, solver='saga', n_jobs=-1, penalty='elasticnet',
↳)
table1.groupby('phenotype')['participant_id'].nunique()
```

```
phenotype
0      113
1      122
Name: participant_id, dtype: int64
```

### 1.8.2 Grid search for 3 hyperparameters

```
# parameters = {'C': [0.005, 0.01, 0.02, 0.05, 0.1, 0.5, 1, 10, 20, 30],
#               'max_iter': [10, 25, 50, 75, 100, 150, 200, 400, 800, 1600],
#               'l1_ratio': [1, 0.9, 0.8, 0.7, 0.6, 0.5, 0.4, 0.3, 0.2, 0.1]}

parameters = {'C': [0.1, 0.5, 1],
               'max_iter': [10, 25],
               'l1_ratio': [0.7, 0.6, 0.5]}

grid_lr = GridSearchCV(lr, parameters, verbose=False, scoring='roc_auc', n_jobs=-1,
↳cv=10)
if not sys.warnoptions:
    warnings.simplefilter("ignore", category=ConvergenceWarning)
    os.environ["PYTHONWARNINGS"] = "ignore" # Also affect subprocesses
grid_lr.fit(X, Y)
```

```
GridSearchCV(cv=10,
             estimator=LogisticRegression(n_jobs=-1, penalty='elasticnet',
                                           random_state=42, solver='saga'),
             n_jobs=-1,
             param_grid={'C': [0.1, 0.5, 1], 'l1_ratio': [0.7, 0.6, 0.5],
                          'max_iter': [10, 25]},
             scoring='roc_auc', verbose=False)
```

### 1.8.3 Best estimator

```
best_lr = grid_lr.best_estimator_

max_auc_score = roc_auc_score(Y, best_lr.predict_proba(X)[:, 1])
coefs = best_lr.coef_[0, :]
num_coef = np.sum(coefs != 0)
X_header = np.array(X.columns)

data_array = np.vstack((X_header, coefs))
model_coefs = pd.DataFrame(data=data_array.T, columns=['SNP', 'Coefficient'])
print(f'Max AUC score: {max_auc_score}\n')
print(f'Non-zero coefficients: {num_coef}\n')
print(f'Best estimator: {grid_lr.best_estimator_}')
print(f'Scorer: {grid_lr.scorer_}')
print(f'Best params: {grid_lr.best_params_}')
print(f'Best score: {grid_lr.best_score_}\n')
m = model_coefs[model_coefs['Coefficient'] != 0].sort_values(by='Coefficient')
m = m.reset_index(drop=True).assign(Index=range(len(m)))
m.Index = m.Index + 1
m.set_index('Index')
```

Max AUC score:0.8581894675758015

Non-zero coefficients: 112

Best estimator: LogisticRegression(C=0.5, l1\_ratio=0.6, max\_iter=10, n\_jobs=-1,  
penalty='elasticnet', random\_state=42, solver='saga')

Scorer: make\_scorer(roc\_auc\_score, needs\_threshold=True)

Best params: {'C': 0.5, 'l1\_ratio': 0.6, 'max\_iter': 10}

Best score: 0.5679924242424242

Index	SNP	Coefficient
1	rs6871718	-0.426813
2	rs999361	-0.426221
3	rs7826007	-0.365302
4	rs3733349	-0.346928
5	rs78197677	-0.284285
6	rs7535292	-0.26294
7	rs4859611	-0.257429
8	rs1949362	-0.254682
9	rs11601088	-0.247447
10	rs7515378	-0.221912
11	rs12101192	-0.193561
12	rs231454	-0.182738
13	rs1859223	-0.170208
14	rs79976845	-0.159289
15	rs4495967	-0.150381
16	rs4482120	-0.150381
17	rs1501467	-0.140416
18	rs79828056	-0.136308
19	rs2433733	-0.12592
20	rs1372420	-0.125209
21	rs41311559	-0.114078
22	rs140859835	-0.084564

(continues on next page)

(continued from previous page)

23	rs11097297	-0.083788
24	rs11097213	-0.076345
25	rs3822023	-0.074309
26	rs142448570	-0.073881
27	rs10448130	-0.073017
28	rs76848738	-0.061152
29	rs11021711	-0.060124
30	rs28691231	-0.049795
31	rs1269243287	-0.049795
32	rs34572188	-0.049509
33	rs2412116	-0.043475
34	rs375017	-0.03843
35	rs7666159	-0.033879
36	rs4416502	-0.030128
37	rs9330264	-0.027804
38	rs734073	-0.023279
39	rs55805734	-0.021777
40	rs8064765	-0.019273
41	rs4796663	-0.015331
42	rs62075803	-0.015331
43	rs12510869	-0.009278
44	rs3850379	-0.000811
45	rs10515816	-0.000433
46	rs3912643	-0.000415
47	rs13188899	-0.00041
48	rs4921799	-0.000409
49	rs4897753	-0.000298
50	rs6964	-0.000209
51	rs3806760	-0.000191
52	rs4028634	-0.000161
53	rs11579790	-0.000147
54	rs379066	-0.000132
55	rs3850745	-0.000123
56	rs4841589	-0.000072
57	rs12499663	-0.000069
58	rs74609071	0.000652
59	rs35456861	0.001887
60	rs9985581	0.009368
61	rs6842271	0.009368
62	rs34288580	0.009444
63	rs148514732	0.012629
64	rs7515370	0.020841
65	rs5848	0.023656
66	rs77312060	0.026649
67	rs72838312	0.02777
68	rs444618	0.029691
69	rs12995314	0.031652
70	rs62190394	0.031652
71	rs10110312	0.033873
72	rs12543164	0.035274
73	rs2075583	0.035907
74	rs6532190	0.039457
75	rs17324625	0.051377
76	rs7682766	0.058691
77	rs143756122	0.060175
78	rs148894916	0.060175

(continues on next page)

(continued from previous page)

79	rs11264302	0.060206
80	rs71371995	0.061233
81	rs12150561	0.063362
82	rs494312	0.066012
83	rs536718528	0.066406
84	rs12935995	0.070272
85	rs2410595	0.070705
86	rs7104332	0.077242
87	rs4331494	0.077355
88	rs7734182	0.085825
89	rs35933728	0.086571
90	rs17016235	0.090649
91	rs34333	0.112067
92	rs201304809	0.120855
93	rs12434297	0.130994
94	rs11589479	0.132404
95	rs9799610	0.162529
96	rs72803476	0.163414
97	rs35776335	0.173641
98	rs2128786	0.189926
99	rs72846765	0.194033
100	rs6599389	0.196567
101	rs73211813	0.203329
102	rs7387252	0.20655
103	rs1800606	0.21257
104	rs208024	0.219757
105	rs9971953	0.254225
106	rs9971789	0.254225
107	rs7520918	0.289942
108	rs13161496	0.291558
109	rs3785891	0.302039
110	rs764324	0.306943
111	rs4241591	0.308725
112	rs11248057	0.32366

## 1.9 NI Females

### 1.9.1 Logistic regression model

```

table1 = table[(table.gender == "M") & (table.inv_genotype=="NI")]
X = table1[table1.columns[5:]]
Y = table1['phenotype']
lr = LogisticRegression(random_state=42, solver='saga', n_jobs=-1, penalty='elasticnet
↳')
table1.groupby('phenotype')['participant_id'].nunique()

```

```

phenotype
0      480
1      477
Name: participant_id, dtype: int64

```



## 1.9.2 Grid search for 3 hyperparameters

```
# parameters = {'C': [0.005, 0.01, 0.02, 0.05, 0.1, 0.5, 1, 10, 20, 30],
#               'max_iter': [10, 25, 50, 75, 100, 150, 200, 400, 800, 1600],
#               'l1_ratio': [1, 0.9, 0.8, 0.7, 0.6, 0.5, 0.4, 0.3, 0.2, 0.1]}

parameters = {'C': [0.005, 0.01],
               'max_iter': [10, 25],
               'l1_ratio': [0.2, 0.1]}

grid_lr = GridSearchCV(lr, parameters, verbose=False, scoring='roc_auc', n_jobs=-1,
cv=10)
if not sys.warnoptions:
    warnings.simplefilter("ignore", category=ConvergenceWarning)
    os.environ["PYTHONWARNINGS"] = "ignore" # Also affect subprocesses
grid_lr.fit(X, Y)
```

```
GridSearchCV(cv=10,
             estimator=LogisticRegression(n_jobs=-1, penalty='elasticnet',
                                           random_state=42, solver='saga'),
             n_jobs=-1,
             param_grid={'C': [0.005, 0.01], 'l1_ratio': [0.2, 0.1],
                          'max_iter': [10, 25]},
             scoring='roc_auc', verbose=False)
```

## 1.9.3 Best estimator

```
best_lr = grid_lr.best_estimator_

max_auc_score = roc_auc_score(Y, best_lr.predict_proba(X)[:, 1])
coefs = best_lr.coef_[0, :]
num_coef = np.sum(coefs != 0)
X_header = np.array(X.columns)

data_array = np.vstack((X_header, coefs))
model_coefs = pd.DataFrame(data=data_array.T, columns=['SNP', 'Coefficient'])
print(f'Max AUC score: {max_auc_score}\n')
print(f'Non-zero coefficients: {num_coef}\n')
print(f'Best estimator: {grid_lr.best_estimator_}')
print(f'Scorer: {grid_lr.scorer_}')
print(f'Best params: {grid_lr.best_params_}')
print(f'Best score: {grid_lr.best_score_}\n')
m = model_coefs[model_coefs['Coefficient'] != 0].sort_values(by='Coefficient')
m = m.reset_index(drop=True).assign(Index=range(len(m)))
m.Index = m.Index + 1
m.set_index('Index')
```

Max AUC score:0.5835888364779874

Non-zero coefficients: 4

Best estimator: LogisticRegression(C=0.005, l1\_ratio=0.1, max\_iter=25, n\_jobs=-1,
penalty='elasticnet', random\_state=42, solver='saga')

(continues on next page)

(continued from previous page)

```

Scorer: make_scorer(roc_auc_score, needs_threshold=True)
Best params: {'C': 0.005, 'l1_ratio': 0.1, 'max_iter': 25}
Best score: 0.5212733636229314

```

Index	SNP	Coefficient
1	rs7520918	-0.011899
2	rs9985581	0.016555
3	rs6842271	0.016555
4	rs3806760	0.027075

## 1.10 II Females

### 1.10.1 Logistic regression model

```

table1 = table[table.gender == "F" & (table.inv_genotype=="II")]
X = table1[table1.columns[5:]]
Y = table1['phenotype']
lr = LogisticRegression(random_state=42, solver='saga', n_jobs=-1, penalty='elasticnet',
↳)
table1.groupby('phenotype')['participant_id'].nunique()

```

```

phenotype
0      195
1      175
Name: participant_id, dtype: int64

```

### 1.10.2 Grid search for 3 hyperparameters

```

# parameters = {'C': [0.005, 0.01, 0.02, 0.05, 0.1, 0.5, 1, 10, 20, 30],
#               'max_iter': [10, 25, 50, 75, 100, 150, 200, 400, 800, 1600],
#               'l1_ratio': [1, 0.9, 0.8, 0.7, 0.6, 0.5, 0.4, 0.3, 0.2, 0.1]}

parameters = {'C': [1, 10, 20],
               'max_iter': [1600, 3200],
               'l1_ratio': [1, 0.9]}

grid_lr = GridSearchCV(lr, parameters, verbose=False, scoring='roc_auc', n_jobs=-1,
↳cv=10)
if not sys.warnoptions:
    warnings.simplefilter("ignore", category=ConvergenceWarning)
    os.environ["PYTHONWARNINGS"] = "ignore" # Also affect subprocesses
grid_lr.fit(X, Y)

```

```

GridSearchCV(cv=10,
             estimator=LogisticRegression(n_jobs=-1, penalty='elasticnet',
                                           random_state=42, solver='saga'),

```

(continues on next page)

(continued from previous page)

```

n_jobs=-1,
param_grid={'C': [1, 10, 20], 'l1_ratio': [1, 0.9],
            'max_iter': [1600, 3200]},
scoring='roc_auc', verbose=False)

```

### 1.10.3 Best estimator

```

best_lr = grid_lr.best_estimator_

max_auc_score = roc_auc_score(Y, best_lr.predict_proba(X)[: , 1])
coefs = best_lr.coef_[0, :]
num_coef = np.sum(coefs != 0)
X_header = np.array(X.columns)

data_array = np.vstack((X_header, coefs))
model_coefs = pd.DataFrame(data=data_array.T, columns=['SNP', 'Coefficient'])
print(f'Max AUC score: {max_auc_score}\n')
print(f'Non-zero coefficients: {num_coef}\n')
print(f'Best estimator: {grid_lr.best_estimator_}')
print(f'Scorer: {grid_lr.scorer_}')
print(f'Best params: {grid_lr.best_params_}')
print(f'Best score: {grid_lr.best_score_}\n')
m = model_coefs[model_coefs['Coefficient'] != 0].sort_values(by='Coefficient')
m = m.reset_index(drop=True).assign(Index=range(len(m)))
m.Index= m.Index + 1
m.set_index('Index')

```

Max AUC score:0.8847765567765569

Non-zero coefficients: 131

Best estimator: LogisticRegression(C=10, l1\_ratio=1, max\_iter=3200, n\_jobs=-1,  
penalty='elasticnet', random\_state=42, solver='saga')  
Scorer: make\_scorer(roc\_auc\_score, needs\_threshold=True)  
Best params: {'C': 10, 'l1\_ratio': 1, 'max\_iter': 3200}  
Best score: 0.5429274165806673

Index	SNP	Coefficient
1	rs138844738	-4.155261
2	rs115879964	-4.155261
3	rs536718528	-3.617586
4	rs115448944	-3.150895
5	rs71607338	-2.57904
6	rs17324625	-2.427067
7	rs2433733	-1.988805
8	rs6842271	-1.715214
9	rs62075803	-1.652955
10	rs4796663	-1.652955
11	rs112957100	-1.549041
12	rs79976845	-1.439649
13	rs74950708	-1.321806

(continues on next page)

(continued from previous page)

14	rs11021711	-1.198849
15	rs1065712	-1.095391
16	rs444618	-1.090401
17	rs13131187	-1.017069
18	rs375017	-1.003348
19	rs11248057	-0.968769
20	rs494312	-0.932152
21	rs3733349	-0.883251
22	rs140859835	-0.836965
23	rs7535292	-0.718728
24	rs7515378	-0.699743
25	rs734073	-0.657171
26	rs76848738	-0.613196
27	rs4482120	-0.610804
28	rs6871718	-0.604013
29	rs112318363	-0.522652
30	rs74609071	-0.505636
31	rs12499663	-0.467277
32	rs77312060	-0.455806
33	rs3744427	-0.437119
34	rs12434297	-0.431097
35	rs11589479	-0.430368
36	rs7666159	-0.427637
37	rs34333	-0.413831
38	rs10110312	-0.399533
39	rs4921799	-0.395144
40	rs1372420	-0.388175
41	rs10515816	-0.384538
42	rs1736103	-0.381249
43	rs73211813	-0.354654
44	rs10448130	-0.33797
45	rs764324	-0.313281
46	rs4841589	-0.298304
47	rs999361	-0.278567
48	rs7387252	-0.254134
49	rs72846765	-0.253638
50	rs71371995	-0.187556
51	rs34992950	-0.183292
52	rs11579790	-0.141663
53	rs9971953	-0.140982
54	rs9971789	-0.140982
55	rs7520918	-0.132976
56	rs35933728	-0.128462
57	rs13161496	-0.126504
58	rs28691231	-0.113809
59	rs4859611	-0.081506
60	rs379066	-0.054031
61	rs72803476	-0.048141
62	rs7826007	-0.021814
63	rs35456861	-0.01947
64	rs148514732	-0.010154
65	rs9330264	-0.008158
66	rs72838312	-0.000916
67	rs12543164	0.007514
68	rs1501467	0.008177
69	rs150107452	0.029482

(continues on next page)

(continued from previous page)

70	rs4241591	0.036761
71	rs1269243287	0.044491
72	rs2128786	0.051102
73	rs2075583	0.066752
74	rs3785891	0.091773
75	rs12995314	0.108912
76	rs62190394	0.108912
77	rs1949362	0.110172
78	rs7682766	0.120521
79	rs6532182	0.12645
80	rs4331494	0.140858
81	rs12935995	0.148705
82	rs1859223	0.150722
83	rs7104332	0.155063
84	rs3850745	0.219926
85	rs11097213	0.221927
86	rs117979807	0.232891
87	rs2410595	0.2445
88	rs35776335	0.287265
89	rs12101192	0.326314
90	rs4495967	0.341328
91	rs79828056	0.353429
92	rs11264302	0.368833
93	rs17016235	0.383153
94	rs6532190	0.383153
95	rs34572188	0.393827
96	rs208024	0.402916
97	rs13188899	0.430538
98	rs6964	0.445559
99	rs75214905	0.458369
100	rs3822023	0.458818
101	rs231454	0.496512
102	rs4897753	0.501664
103	rs11097297	0.514293
104	rs148894916	0.565314
105	rs143756122	0.565314
106	rs201304809	0.586162
107	rs10831599	0.612322
108	rs7515370	0.632874
109	rs9799610	0.642287
110	rs2412116	0.660505
111	rs112270735	0.673886
112	rs12150561	0.82272
113	rs7734182	0.898797
114	rs6836715	0.936025
115	rs78197677	0.968978
116	rs6599389	1.057432
117	rs34288580	1.086857
118	rs5848	1.09313
119	rs3806760	1.103029
120	rs142448570	1.167941
121	rs41311559	1.289618
122	rs9985581	1.394087
123	rs3850379	1.404536
124	rs55805734	1.418262
125	rs117230705	1.559671

(continues on next page)

(continued from previous page)

126	rs74677851	1.919257
127	rs4028634	2.386557
128	rs11601088	2.926625
129	rs1800606	3.409027
130	rs8064765	5.881005
131	rs79531911	7.603942